
Neural Green’s Functions

Seungwoo Yoo Kyeongmin Yeo Jisung Hwang Minhyuk Sung
KAIST
`{dreamy1534, aaaaa, 4011hjs, mhsung}@kaist.ac.kr`

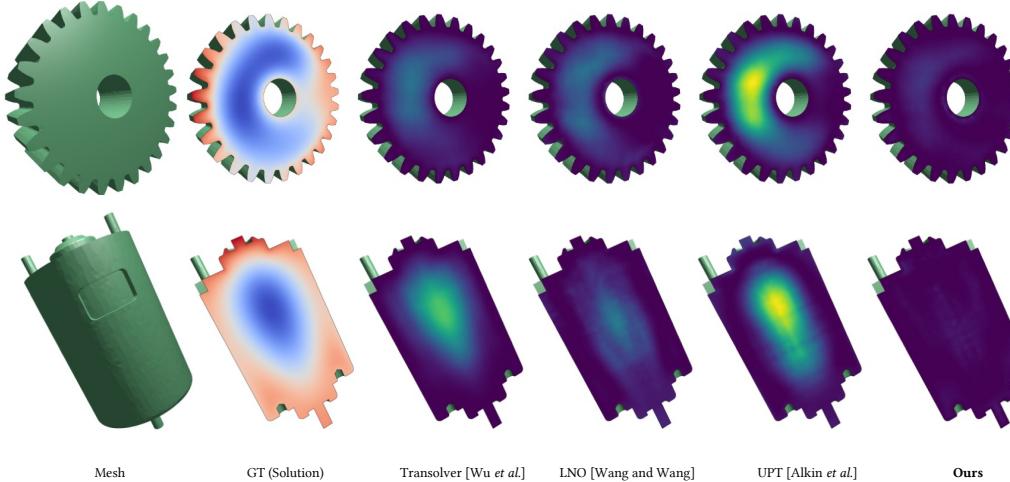


Figure 1: Neural Green’s Function. We propose a neural solution operator designed for linear PDEs whose differential operators admit eigendecompositions. Our framework effectively handles irregular geometries and diverse source and boundary functions. In the steady-state thermal analysis used as a model problem, our method demonstrates superior generalization performance compared to state-of-the-art neural operators in predicting ground-truth solutions (column 2), as evidenced by the error maps (columns 3–6).

Abstract

We introduce Neural Green’s Function, a neural solution operator for linear partial differential equations (PDEs) whose differential operators admit eigendecompositions. Inspired by Green’s functions, the solution operators of linear PDEs that depend exclusively on the domain geometry, we design Neural Green’s Function to imitate their behavior, achieving superior generalization across diverse irregular geometries and source and boundary functions. Specifically, Neural Green’s Function extracts per-point features from a volumetric point cloud representing the problem domain and uses them to predict a decomposition of the solution operator, which is subsequently applied to evaluate solutions via numerical integration. Unlike recent learning-based solution operators, which often struggle to generalize to unseen source or boundary functions, our framework is, by design, agnostic to the specific functions used during training, enabling robust and efficient generalization. In the steady-state thermal analysis of mechanical part geometries from the MCB dataset, Neural Green’s Function outperforms state-of-the-art neural operators, achieving an average error reduction of 13.9% across five shape categories, while being up to 350 times faster than a numerical solver that requires computationally expensive meshing.

1 Introduction

Linear partial differential equations (PDEs) play a central role in many scientific and engineering disciplines, such as thermal analysis, electrostatics, fluid dynamics, and elasticity. Similar to other PDEs, their ubiquity has driven the development of numerical methods, such as finite difference methods (FDMs) and finite element methods (FEMs), to solve the equation defined over complex geometric domains in real-world applications where analytical solutions are often unavailable. These techniques discretize problem domains into fine-grained meshes of lattices or polygons, resulting in linear systems that can be constructed and solved efficiently. However, the reliance on meshes poses a significant limitation, as generating them requires running computationally expensive meshing algorithms [13, 14], making it challenging to rapidly evaluate solutions across multiple problem domains with different boundaries. This bottleneck is especially pronounced during the early design phase of engineering workflows, where rapid iteration is essential for achieving optimal results.

Recently, learning-based solvers have emerged as promising surrogates to conventional numerical solvers for solving PDEs. They predict solutions without mesh construction, either by optimizing a neural network to parameterize the solution function for a given problem instance [30, 19] or by learning solution operators that map input functions to solutions, enabling inference with a single forward pass [20–22, 26, 23, 12, 36, 35, 34, 3]. While these approaches significantly improve efficiency compared to numerical solvers and provide less expensive approximations, little attention has been paid to their generalization capabilities under simultaneous variations in problem domains, source functions, and boundary functions. Some previous methods focus on learning the solutions to PDE problems within a single domain [30, 25, 27], requiring re-training whenever the problem domain or even a specific problem instance changes. Approaches such as Boullé et al. [6], Teng et al. [31], and Negi et al. [28] address this by learning the solution operator for a given domain, enabling the handling of different problem instances. However, they still require re-training when applied to different domains. Others [20–22, 26, 23, 12, 36, 35] attempt to learn solution operators that generalize across a collection of shapes but often struggle to handle unseen source functions and boundary functions. This limitation arises because these inputs are directly provided to the neural networks that regress the solution operators, requiring sufficient training examples to achieve generalization.

In this work, we propose a learnable solution operator capable of handling various shapes and functions, with a specific focus on linear PDEs whose differential operators admit eigendecompositions, such as the Poisson and Biharmonic equations. The foundation of our approach is rooted in the mathematical definition of the solution operator for a linear PDE, known as the *Green’s function*, which only depends on the geometry of the problem domain. Once the Green’s function of a domain is known, the linearity allows solutions to be computed for arbitrary source and boundary functions. This unique property motivated the design of our framework, which predicts neural features solely from the domain geometry. These features are used to approximate the eigendecomposition of the corresponding Green’s function and to predict related differential quantities, such as per-vertex masses, that are essential for computing solutions via numerical integration. This design imposes a strong prior on the solution operator learned by the neural network, making the model independent of the specific source and boundary functions used during training.

In our experiments, we first empirically validate the generalizability of the proposed framework using simple examples of the Poisson and Biharmonic equations. We then extend the evaluation to a practical setting of steady-state thermal analysis on complex 3D geometries, using a benchmark we built from the MCB dataset [16]. This benchmark comprises a diverse collection of mechanical part geometries, each paired with a steady-state thermal distribution obtained by solving Poisson’s equation. The dataset features a wide variety of shapes with significant intra-class variations, along with triplets of source, boundary, and solution functions generated by a numerical solver, providing a challenging benchmark setup for both baselines and our method. In comparisons, Neural Green’s Function demonstrates superior generalizability compared to state-of-the-art neural operators [35, 34, 3]. Notably, our framework reduces the error metric by 13.9 % compared to Transolver [35], while sharing the same backbone. These results highlight the significance of explicitly formulating and predicting solution operators in achieving superior generalization capabilities.

2 Related Work

Physics-Informed Neural Networks (PINNs). Physics-Informed Neural Networks (PINNs) [30] parameterize the solution of each PDE instance directly using a neural network, and optimize its

parameters by minimizing objective functions derived from the governing equations and empirical observations. They can be easily implemented using automatic differentiation, which is supported by deep learning frameworks [29, 2, 7]. However, PINNs are trained separately for each problem instance, with each neural network approximating a single solution at a time. This approach requires retraining whenever the problem domain, source function, or boundary condition changes.

Neural Operators. Unlike PINNs, Neural Operators learn function-to-function mappings that act as solution operators for PDEs. They achieve this by parameterizing (nonlinear) kernel functions and implicitly perform kernel integration across subsequent network layers. Specifically, GNO [20] employs graph neural networks (GNNs) to approximate integral kernels through local aggregation via message-passing. FNO [21] improves efficiency and performance by utilizing fast Fourier transform (FFT) [10] in regular domains and learning global integral kernels in the spectral domain. Follow-up work [22, 26] combine GNO [20] and FNO [21] to handle irregular problem domains by efficiently modeling both local and global interactions within systems. Several work [24, 23, 12, 36] utilizing Transformers [33] architecture have also been proposed, leveraging self-attention mechanism to capture interactions among mesh points. These work bypass the quadratic complexity of self-attention layers using linear Transformers [18, 9, 8]. Recent work further reduce the computational complexity by incorporating compact latent representations with dimensionalities significantly smaller than mesh sizes [35, 34, 3]. Neural operators are a versatile framework for operator learning over irregular geometries with diverse source and boundary functions. However, they often couple input meshes with sampled function values, which makes generalization to new functions challenging.

Learning Green’s Functions. Several recent work focus on linear PDEs and propose data-driven approaches to discovering the Green’s function of a problem domain, enabling solutions for varying source and boundary functions. Boullié et al. [6] propose a method to learn Green’s functions for linear PDEs using pairs of forcing and solution functions, leveraging rational neural networks [5]. Teng et al. [31] extend this approach by enabling the evaluation of solutions with arbitrary source and boundary functions, regressing Green’s functions through the approximation of Dirac delta functions with Gaussian distributions. Similarly, Negi et al. [28] approximate free-space Green’s functions using a radial basis function (RBF) kernel-based neural network. By directly learning Green’s functions, the aforementioned approaches can predict solutions for PDEs with varying source and boundary functions. However, these methods are restricted to individual problem domains and require retraining for unseen geometries, as the networks learn domain-specific Green’s functions and their gradients. Additionally, they focus on simple domains (e.g., circular domains in 2D), where numerical quadrature for computing solutions with predicted Green’s functions is relatively straightforward. Our work addresses these limitations through a framework design inspired by the eigendecomposition of Green’s functions for linear PDEs. In particular, the framework learns to extract geometric features from problem domains and composes them to reconstruct the corresponding Green’s functions. Moreover, our approach predicts differential quantities needed to approximate numerical integration over complex, irregular geometries, extending its applicability to real-world scenarios, as demonstrated in Sec. 5.

3 Background

Consider a linear PDE subject to Dirichlet boundary conditions defined over a continuous domain $D \subset \mathbb{R}^d$ with boundary ∂D :

$$\begin{cases} \mathcal{L}u(x) = f(x) & x \in D \\ u(x) = h(x) & x \in \partial D, \end{cases} \quad (1)$$

where \mathcal{L} is a linear differential operator, and u , f , and h are functions that reside in Banach spaces \mathcal{U} , \mathcal{F} , and \mathcal{H} of functions, respectively. The solution operator \mathcal{G} for Eqn. 1 that computes the solution $u = \mathcal{G}(D, f, h)$ is the mapping:

$$\mathcal{G} : \mathcal{D} \times \mathcal{F} \times \mathcal{H} \rightarrow \mathcal{U}, \quad (2)$$

where \mathcal{D} is a set of problem domains. Since Eqn. 1 is a linear PDE, \mathcal{G} is explicitly given as:

$$\begin{aligned} u(x) &= \mathcal{G}_D(f, h) \\ &= \int_D G_D(x, y) f(y) dy \\ &\quad + \int_{\partial D} h(y) \nabla_y G_D(x, y) \cdot n(y) dy, \end{aligned} \quad (3)$$

where $\mathcal{G}_D := \mathcal{G}(D, \cdot, \cdot)$ is an instantiation of \mathcal{G} in the domain D and $n(y)$ is the outward normal vector at $y \in \partial D$, respectively. The integral kernel $G_D : D \times D \rightarrow \mathbb{R}$, referred to as the *Green’s function* of the operator \mathcal{L} in D , represents its impulse response and is a solution of:

$$\begin{cases} \mathcal{L}G_D(x, y) = \delta(x - y), & x \in D \\ G_D(x, y) = 0, & x \in \partial D \end{cases} \quad (4)$$

where δ is the Dirac delta function. One important property of \mathcal{G}_D and Green’s function G_D is that they only depend on the geometry of the domain D , regardless of the specific source function f or boundary function h . This implies that, once the Green’s function G_D for a given domain D is known, the PDE in Eqn. 1 can be solved for arbitrary source and boundary functions f and h . However, solutions based on Eqn. 3 have only been applied to simple domains where closed-form expressions for Green’s functions are available.

Numerical solvers, on the other hand, utilize a volumetric mesh representing D to discretize the operator \mathcal{L} in Eqn. 1, forming a discrete counterpart represented as linear systems. Assume that D is represented as a mesh $\mathbf{D} = (\mathbf{V}, \mathbf{T})$ of a single connected component, consisting of vertices \mathbf{V} and faces \mathbf{T} . We define a selection matrix $\mathbf{S} \in \{0, 1\}^{N_b \times N_v}$ to indicate N_b boundary vertices among the N_v vertices, and its complement $\mathbf{K} \in \{0, 1\}^{(N_v - N_b) \times N_v}$ to represent the remaining interior vertices. The scalar-valued source function f and the boundary function h are discretized by sampling their values at the vertices, resulting in two vectors: $\mathbf{f} \in \mathbb{R}^{N_v}$ and $\mathbf{h} \in \mathbb{R}^{N_b}$, respectively. Then, Eqn. 1 can be written as a linear system in its discrete form:

$$\mathbf{Lu} = \mathbf{Mf} \quad \text{s.t.} \quad \mathbf{Su} = \mathbf{h}, \quad (5)$$

where $\mathbf{L} \in \mathbb{R}^{N_v \times N_v}$ is the discretization of \mathcal{L} and $\mathbf{M} \in \mathbb{R}^{N_v \times N_v}$ is the mass matrix of the mesh \mathbf{D} , respectively. The solution to this system can be written as:

$$\begin{aligned} \mathbf{u} &= \mathbf{K}^T \left\{ (\mathbf{KLK}^T)^{-1} (\mathbf{KMf} - \mathbf{KLS}^T \mathbf{h}) \right\} + \mathbf{S}^T \mathbf{h} \\ &= \mathbf{K}^T \left\{ \mathbf{G} (\mathbf{KMf} - \mathbf{KLS}^T \mathbf{h}) \right\} + \mathbf{S}^T \mathbf{h}, \end{aligned} \quad (6)$$

where $\mathbf{G} = (\mathbf{KLK}^T)^{-1} \in \mathbb{R}^{(N_v - N_b) \times (N_v - N_b)}$ represents the inverse of the matrix \mathbf{L} restricted to the interior of the domain. The detailed derivation can be found in the Appendix. The matrix \mathbf{G} serves the same role as the Green’s function discussed in Sec. 3. Notably, the terms \mathbf{G} , \mathbf{M} , and \mathbf{K} in Eqn. 6 are independent of the given functions \mathbf{f} and \mathbf{h} , analogous to the Green’s function G_D in Eqn. 3. This independence implies that the solution operator for Eqn. 5 can be approximated by predicting these terms using only the geometric information of the mesh \mathbf{D} .

Assuming that \mathbf{L} is symmetric, its submatrix \mathbf{KLK}^T , obtained by eliminating the rows and columns corresponding to the boundary vertices, is also symmetric. Furthermore, if the imposed boundary conditions are sufficient, \mathbf{KLK}^T becomes full-rank and admits the following eigendecomposition:

$$\mathbf{G} = \Phi \Lambda^{-1} \Phi^T, \quad (7)$$

where $\Phi \in \mathbb{R}^{(N_v - N_b) \times (N_v - N_b)}$ is the matrix whose columns are the eigenvectors of \mathbf{KLK}^T , and $\Lambda \in \mathbb{R}^{(N_v - N_b) \times (N_v - N_b)}$ is the diagonal matrix of positive eigenvalues corresponding to the eigenvectors in Φ . These properties form the basis of our novel framework, designed to predict solutions of linear PDEs across diverse irregular geometries, as detailed in the following section.

4 Neural Green’s Function

Numerical solvers are well-established tools for solving linear PDEs; however, they rely on the volumetric mesh \mathbf{D} generated from surface representations (e.g., boundary representations in CAD), which can be computationally expensive. This poses a challenge to rapid iteration, which is essential in the early stages of design exploration.

Directly applying neural surrogate models [20–22, 26, 24, 23, 12, 36, 35] may help mitigate this issue by enabling coordinate-based solution queries, eliminating the need for the mesh \mathbf{D} . These models effectively capture complex behaviors of nonlinear systems by encoding both query points, along with source and boundary functions, into learned latent representations. However, for the class of linear PDEs discussed in Sec. 3, we claim that a stronger prior can be incorporated into the

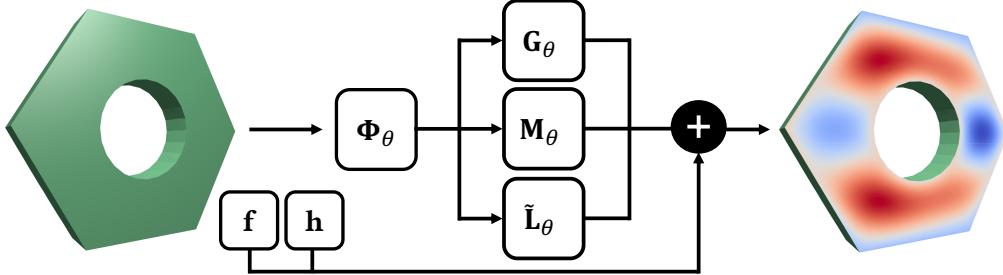


Figure 2: Method overview. Given query points representing the geometry of the problem domain where the solution function is to be predicted, Neural Green’s Function extracts neural features that are subsequently used to approximate the Green’s function and to predict the differential quantities required to evaluate Eqn. 6. The example is rendered using a tetrahedral mesh, which is used solely for visualization purposes.

operator design to improve its generalization capability. In the following, we discuss how the design of Neural Green’s Function, inspired by Green’s functions, enables generalization to unseen source and boundary functions. Although we mainly focus on problems in 3D space, our approach is not limited to a particular dimension.

Our framework, as outlined in Fig. 2, takes a set of query points \mathbf{Q} in 3D space that represents the geometry of the domain D . These points can be obtained through interior point queries [4], thereby avoiding the costly domain meshing. From these query points, we extract features that are used to approximate Green’s functions at the corresponding points, rather than directly predicting the solutions. While the query points are associated with a source function f evaluated at \mathbf{Q} and a boundary function h sampled at the boundary points, our network remains agnostic to these functions. This independence is the key for achieving generalization to unseen functions.

Specifically, we employ a neural network with parameters θ to predict features $\Phi_\theta \in \mathbb{R}^{|\mathbf{Q}| \times d}$ for all query points, whose coordinates are provided as inputs to the network. These predicted features are then utilized to construct our neural Green’s function \mathbf{G}_θ , which serves as a drop-in replacement for the ground-truth Green’s function \mathbf{G} in Eqn. 6. Inspired by Eqn. 7, we compute each element of \mathbf{G}_θ as the dot product of feature vectors predicted for two interior points. The computation of \mathbf{G}_θ for all pairs of points is efficiently parallelized through a single matrix multiplication:

$$\mathbf{G}_\theta = (\mathbf{K}\Phi_\theta)(\mathbf{K}\Phi_\theta)^T. \quad (8)$$

Since the mass matrix \mathbf{M} and the submatrix \mathbf{KLS}^T of the operator \mathbf{L} are unavailable without a mesh \mathbf{D} , we also predict these quantities to evaluate Eqn. 6. In particular, we employ an MLP to decode per-vertex features Φ_θ into per-vertex mass values \mathbf{M}_θ , using the Softplus activation to ensure that the predicted masses are positive. Similarly to \mathbf{G}_θ , the submatrix $\tilde{\mathbf{L}} = \mathbf{KLS}^T$ is predicted by computing dot products between features $\Psi_\theta \in \mathbb{R}^{|\mathbf{Q}| \times d}$, which are decoded from Φ_θ by another MLP:

$$\tilde{\mathbf{L}}_\theta = (\mathbf{K}\Psi_\theta)(\mathbf{S}\Psi_\theta)^T. \quad (9)$$

Lastly, these components are combined to evaluate Eqn. 6:

$$\mathbf{u}_\theta = \mathbf{K}^T \left\{ \mathbf{G}_\theta \left(\mathbf{K}\mathbf{M}_\theta \mathbf{f} - \tilde{\mathbf{L}}_\theta \mathbf{h} \right) \right\} + \mathbf{S}^T \mathbf{h}, \quad (10)$$

yielding the predicted solution \mathbf{u}_θ .

Given a dataset of N examples $\{(\mathbf{u}^i, \mathbf{f}^i, \mathbf{h}^i)\}_{i=1}^N$ generated by solving Eqn. 5 using an FEM solver on meshes $\{\mathbf{D}^i\}_{i=1}^N$ along with their corresponding mass matrices $\{\mathbf{M}^i\}_{i=1}^N$, the network parameters θ are optimized end-to-end by minimizing the empirical risk:

$$\hat{\theta} = \arg \min_{\theta} \mathbb{E}_i [\|\mathbf{u}^i - \mathbf{u}_\theta^i\|_2^2 + \lambda \|\mathbf{M}_\theta^i - \text{diag}(\mathbf{M}^i)\|_2^2], \quad (11)$$

where $\lambda \geq 0$ is a regularization weight applied to the predicted mass values to ensure consistency with the ground truth. We observe that regularizing the predicted masses is essential for the convergence of network training, as discussed further in Sec. 5.4. In all our experiments, we set $\lambda = 1$.

5 Experiment

5.1 Experiment Setup

In this section, we outline the baselines, evaluation metric, and implementation details of our experiments, while deferring detailed descriptions of the problem setups to their respective sections.

Baselines. For 2D Poisson and Biharmonic examples in Sec. 5.2, we compare Neural Green’s Function against Transolver [35], the state-of-the-art neural operator. In Sec. 5.3, we further expand the set of baselines to conduct a more comprehensive evaluation in steady-state thermal analysis. This includes Transolver [35], Latent Neural Operator (LNO) [34], and Universal Physics Transformer (UPT) [3]. For all baselines, we use their official implementations and train the networks under the same configuration as ours. For neural operators that require conditioning on source and boundary functions, we follow the approach of Transolver [35] and LNO [34], concatenating the source and boundary function values at the query points as inputs to the baseline models.

Evaluation Metrics. Following the baselines [35, 34, 3], the error $e(\mathbf{u}, \mathbf{u}_\theta)$ between the predicted solution \mathbf{u}_θ and the corresponding ground truth \mathbf{u} is measured as the relative L_2 distance:

$$e(\mathbf{u}, \mathbf{u}_\theta) = \frac{\|\mathbf{u}_\theta - \mathbf{u}\|_2}{\|\mathbf{u}\|_2}, \quad (12)$$

which is averaged over the entire test set for quantitative comparisons.

Implementation Details. We adopt the network architecture of Transolver [35], a state-of-the-art neural operator and one of our baseline methods, to extract the features Φ_θ , encouraging feature sharing across query points. Specifically, our network consists of an MLP that maps coordinates of query points to latent representations. The layer is followed by eight Transolver blocks [35] and MLP heads that decode them into the components required to evaluate Eqn. 10. Unless otherwise specified, all models in our experiments are trained for 40 epochs using the ADAM optimizer [17] with a OneCycleLR learning rate scheduler, setting the maximum learning rate to 1×10^{-4} . For the steady-state thermal analysis experiments in Sec. 5.3, we use a batch size of 1 and accumulate gradients over 8 training steps to stabilize training, to handle meshes with a large number of vertices.

5.2 Two-Dimensional Poisson and Biharmonic Equations

We first assess the generalization capability of Neural Green’s Function to unseen source and boundary functions, by fixing the problem domain to a unit square $[0, 1] \times [0, 1]$ in 2D. The domain is discretized with the resolution of 100×100 by uniformly placing grid points along the two axes. As model problems, we consider Poisson and Biharmonic equations defined over this domain.

Problem instances of Poisson’s equation are generated using boundary functions instantiated from the template that satisfies $\Delta u = 0$:

$$u(x, y) = A(x^3 - 3xy^2) + B(y^3 - 3x^2y) + x^2, \quad (13)$$

with the source function set to zero within the domain interior. We use 100 training and 100 test examples, each generated by sampling tuples of coefficients (A, B) . To ensure that the boundary functions at test time do not overlap with those used for training, we sample coefficients from the uniform distribution $\mathcal{U}[-1, 1]$ for the training set, and from $\mathcal{U}[1, 2]$ for the test set.

For Biharmonic equation, we consider the template

$$u(x, y) = A(x^4 - 6x^2y^2 + y^4) + Bx^4 + Cy^4, \quad (14)$$

which satisfies $\Delta^2 u = 0$. Analogous to the Poisson’s equation setup, we use 100 training and 100 test examples, generated by sampling coefficients over different ranges. For the training set, the coefficients (A, B, C) are drawn from $\mathcal{U}[-1, 1]$, while for the test set, they are sampled from $\mathcal{U}[1, 2]$.

For both setups, we train our Neural Green’s Function and Transolver [35] for 40 and 200 epochs, respectively. We increase the number of training epochs for Transolver to account for its slower convergence in training loss. After training, we evaluate both models on the test set by computing the relative L_2 error across all examples. As summarized in Tab. 1, Neural Green’s Function, which

	Poisson's Equation		Biharmonic Equation	
	Train Set	Test Set	Train Set	Test Set
Transolver (Wu et al. [35])	0.053	0.372	0.025	0.337
Ours	0.014	0.012	0.010	0.009

Table 1: Relative L_2 errors measured on the training and test sets of 2D Poisson and Biharmonic equation examples. In both setups, our method achieves comparable errors on the training and test sets, whereas Transolver [35] exhibits a notable gap due to its network being jointly conditioned on both domain geometry and boundary functions.

	SCREWS & BOLTS	NUT	MOTOR	FITTING	GEAR
Transolver (Wu et al. [35])	<u>0.221</u>	0.320	0.407	0.180	0.281
LNO (Wang & Wang [34])	0.239	0.372	0.528	0.259	0.466
UPT (Alkin et al. [3])	0.358	0.516	0.765	0.392	0.507
Ours	0.189	0.275	0.338	0.160	0.243
Error Reduction (%)	14.7	14.1	16.9	10.8	13.3

Table 2: Relative L_2 errors measured on the test sets of the steady-state thermal analysis benchmark. Our method achieves the lowest relative L_2 error across all five shape categories. Notably, it improves the metric of Transolver [35] by an average of 13.9%, despite utilizing the same network architecture. The error reduction is the difference between the two errors, divided by the second-best error.

relies solely on features extracted from the domain geometry while remaining agnostic to the specific boundary functions used during training, generalizes well compared to Transolver [35], whose network is jointly conditioned on both domain geometry and boundary functions.

5.3 Steady-State Thermal Analysis

Building on the observations from Sec. 5.2, we extend our experiments to a more challenging and practical task, steady-state thermal analysis on complex 3D geometries. This task is appropriate for evaluating the generalizability of our framework to unseen problem domains, beyond the source and boundary functions considered in the previous section. To this end, we construct a new PDE benchmark using the MCB dataset [16], which contains a variety of 3D mechanical part shapes. We generate tetrahedral meshes for shapes in five categories (SCREWS & BOLTS, NUT, MOTOR, FITTING, and GEAR) by meshing the interiors of unit-cube-normalized shapes using fTetWild [14]. The shape collection is divided into 200 shapes for training and 20 shapes for testing to evaluate generalization to unseen problem domains. We illustrate several example shapes in Fig. 3. As shown, the shape collection used in the experiments includes shapes with thin structures and intricate details. Additionally, shapes within the same category can have significantly different geometries, presenting a challenge for learned solution operators to generalize effectively across diverse geometries.

To generate PDE examples, we employ an off-the-shelf FEM solver [11] to solve Poisson's equation defined on the shapes in our dataset. These problems are constructed using different source and boundary functions derived from predefined templates, details of which are provided in the Appendix. Specifically, we use 8 source functions for training and reserve an additional 8 for testing. For boundary functions, we use 2 for training and 2 for testing. This setup results in 16 unique combinations of source and boundary functions for training, while 16 entirely unseen combinations are used for testing. It offers greater problem diversity than other benchmarks for PDEs on irregular geometries, such as ShapeNet-CFD [32], which assumes a fixed driving velocity. Overall, the dataset comprises 3200 shape-problem pairs for training and 320 pairs for testing across all categories.



Figure 3: Example shapes from our dataset. Our dataset comprises diverse mechanical part shapes from the MCB dataset [16], designed to evaluate the generalizability of learning-based PDE solvers across shape variations.

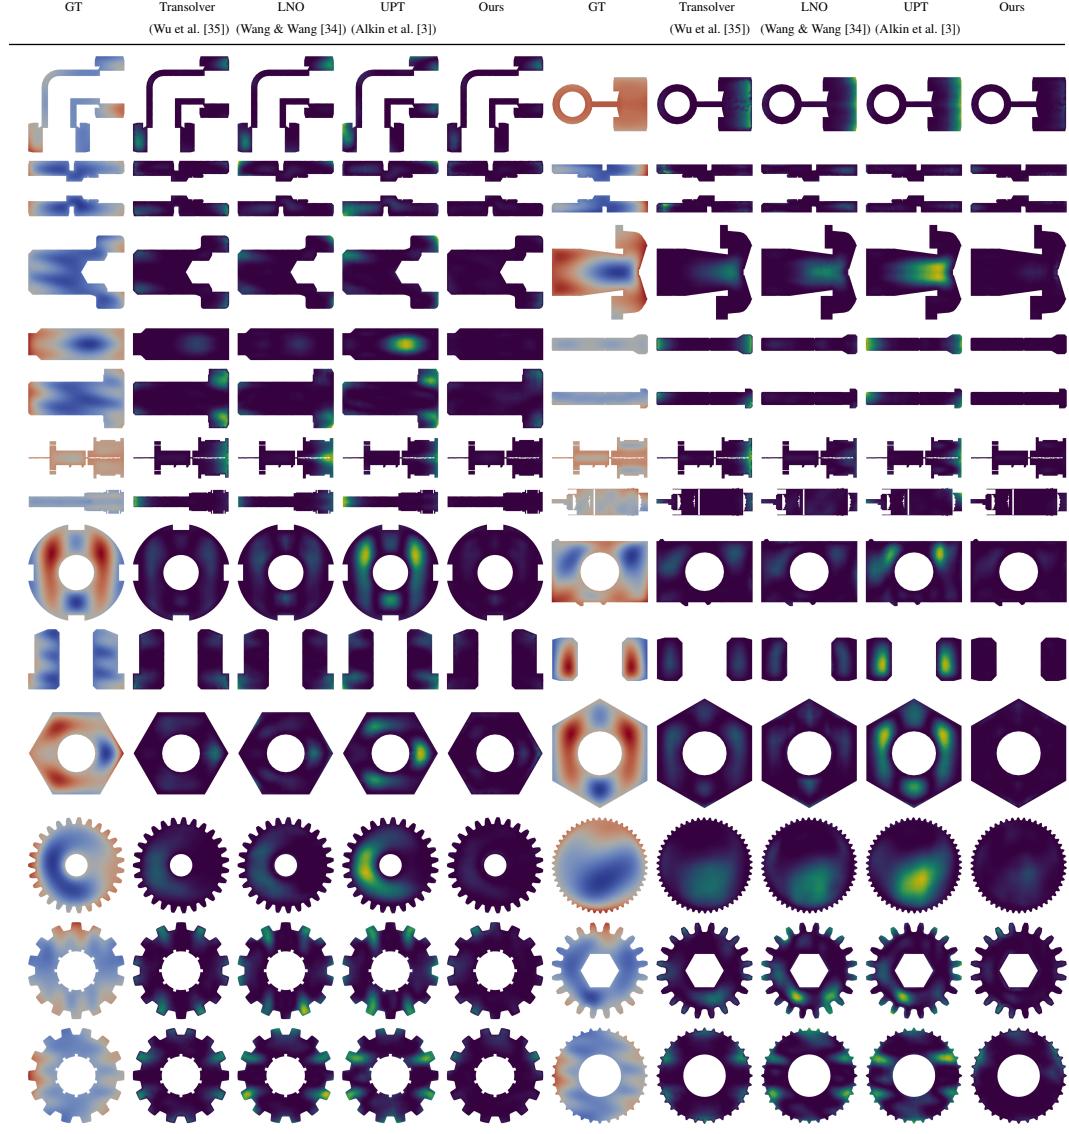


Figure 4: Qualitative comparison. For each ground-truth solution in columns one and six, the per-point L_2 errors from various methods are visualized in the error maps across the remaining columns. Red denotes higher values in the columns labeled “GT”, while in the error maps, brighter colors indicate greater errors. Neural Green’s Function demonstrates its ability to accurately predict solutions on irregular geometries from diverse domains. Compared to state-of-the-art neural operators trained to directly map source and boundary functions to solutions, our method achieves lower errors, as indicated by the darker regions in the error maps. The examples are rendered using tetrahedral meshes, that are used solely for visualization purposes. Best viewed when zoomed-in.

Quantitative Analysis. In Tab. 2, we summarize the errors measured on the test sets of five shape categories using the baselines and our method. Neural Green’s Function consistently achieves lower errors compared to the baselines, demonstrating the effectiveness of the proposed method that explicitly predicts the solution operator. Compared to Transolver [35] which achieves the second-best errors across all categories, our method achieves an average error reduction of 13.9%, calculated as the difference between the best and second-best errors, divided by the latter. Since our framework utilizes the network architecture of Transolver [35] as its backbone, this improvement highlights the significance of incorporating a prior on the solution operator.

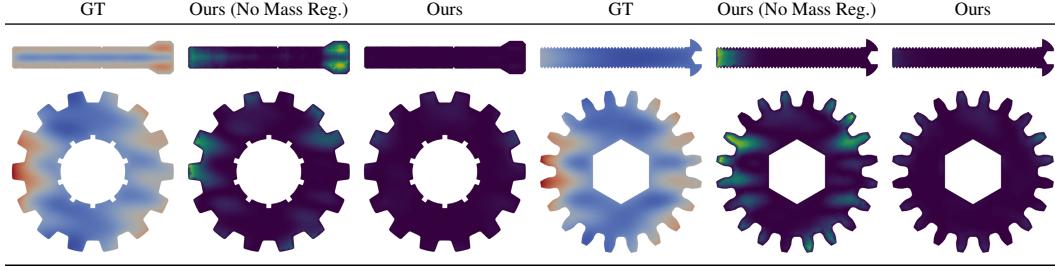


Figure 5: Qualitative results from the ablation study. For each ground-truth solution in columns one and four, we visualize the per-point L_2 errors from variants of our method in the remaining columns. Red denotes higher values in the columns labeled “GT”, while in the error maps, brighter colors indicate greater errors. The examples are rendered using tetrahedral meshes, that are used solely for visualization purposes. Best viewed when zoomed-in.

Qualitative Analysis. Qualitative results from five categories—FITTING (rows 1–2), SCREWS & BOLTS (rows 3–5), MOTOR (rows 6–7), NUT (rows 8–10), and GEAR (rows 11–13)—are presented in Fig. 4. Additional results can be found in the Appendix. For each example, we display color maps of the ground-truth solutions from the dataset, labeled as “GT”, alongside error maps of the predictions generated by different methods. In the solution visualizations, red indicates higher values, while in the error maps, brighter colors represent greater errors. We visualize the slices taken along one of the canonical axes (x , y , z) of each shape for clarity. The darker error maps, compared to those of the baselines, demonstrate that our method produces more accurate predictions, validating our observations from the quantitative analysis. Our method better handles irregular shapes with intricate details, such as thin structures of MOTOR shapes (rows 6–7), or the cogs of GEAR shapes (rows 11–13).

Runtime Analysis. We evaluate the efficiency of our framework by comparing it with an FEM solver and neural operator baselines [35, 34, 3], measuring the total runtimes required to process the test sets. The runtimes are summarized in Tab. 6. All reported runtimes are measured in seconds, with each sample processed sequentially to minimize system load and prevent interference that could impact the timing of other samples. Compared to the FEM solver, Neural Green’s Function benefits greatly from not requiring mesh generation, which accounts for the majority of the processing time in the FEM pipeline. Notably, it achieves a substantial reduction in runtime, being up to 350 times faster than FEM. Furthermore, Neural Green’s Function introduces only minimal computational overhead compared to neural operator baselines [35, 34, 3] while delivering superior generalizability and performance, as discussed previously. Details of this analysis, including the breakdown of the runtimes, are provided in the Appendix.

	FEM Total (s)	Transolver (s) (Wu et al. [35])	LNO (s) (Wang & Wang [34])	UPT (s) (Alkin et al. [3])	Ours (s)
SCREWS & BOLTS	12.956	0.033	0.022	0.056	0.039
NUT	12.238	0.022	0.020	0.076	0.051
MOTOR	50.095	0.090	0.088	0.166	0.140
FITTING	18.439	0.032	0.030	0.076	0.059
GEAR	46.384	0.188	0.187	0.246	0.225

Table 3: Runtime comparison against an FEM solver and neural operators. Our approach achieves up to a $350\times$ speedup over the FEM solver, which requires mesh generation, while maintaining comparable runtime to neural operator baselines and delivering improved generalizability and performance. All runtimes are reported in seconds.

5.4 Ablation Study

We analyze the impact of mass regularization during training and the effect of feature dimensionality on model performance. In Tab. 4, we report the test errors of our model and its variant trained without mass regularization, denoted as “Ours (No Mass Reg.)”, on the SCREWS & BOLTS and GEAR categories. As reflected in the errors, regularizing the mass prediction outputs from our model

	SCREWS & BOLTS	GEAR
Ours (No Mass Reg.)	0.285	0.411
Ours	0.189	0.243

Table 4: Quantitative results from the ablation study. Regularization of the mass predictions is crucial for performance.

d	64	128	256
Relative L_2	0.180	0.189	0.206

Table 5: Analysis of feature dimension. The model’s performance remains consistent across different feature dimension choices.

is crucial for performance. This is because the per-vertex masses involved in Eqn. 10 are typically very small (approximately on the order of 1×10^{-4}), leading to instability during the early stages of network training when the initial mass predictions deviate significantly in scale from the ground-truth values. Qualitative results showcased in Fig. 5 further support this observation: the model trained with mass regularization produces more accurate predictions, as evidenced by the darker regions in the error maps (columns 2–3 and 5–6).

On the other hand, we examine whether the model’s performance is sensitive to the dimensionality d of the feature representation Φ_θ . Intuitively, Φ_θ plays a role of the eigenvectors in Eqn. 7, serving as low-rank bases for approximating the operator matrix. To assess this, we train variants of our framework, which by default uses 128-dimensional feature vectors, with alternative configurations employing 64 and 256 dimensions. All models are trained on the SCREWS & BOLTS class using the same setup as in our main experiments, and the relative L_2 errors on the test set are reported in Tab. 5. As shown, the model’s performance remains insensitive to the choice of feature dimension, highlighting that even 64 learned bases are sufficient to accurately approximate the solution operator.

6 Conclusion

We present Neural Green’s Function, a solution operator specialized in linear PDEs whose differential operators admit eigendecompositions, capable of generalizing across diverse domains, source and boundary functions. Our framework incorporates the principles of Green’s functions into its design by extracting neural features solely from the geometries of problem domains. These features are then used to approximate the corresponding Green’s functions and related differential quantities, which are subsequently utilized to perform numerical integration for solution prediction. We empirically validate our design choice by applying Neural Green’s Function to solve the Poisson and Biharmonic equations on 2D domains, and further demonstrate that it outperforms state-of-the-art neural operators on a challenging 3D benchmark encompassing a wide variety of shapes and Poisson equation instances defined over them. In particular, Neural Green’s Function achieves a 13.9% improvement over the best performing baseline, while accelerating solution evaluation by $350\times$ in steady-state thermal analysis involving complex 3D geometries. These results highlight the superior generalizability of our framework across problem domains, as well as diverse source and boundary functions.

Limitations and Future Work. While we believe our approach represents a step toward developing generalizable, data-driven solution operators for PDEs, it is not without limitations. In particular, extending our framework to a broader class of linear PDEs and incorporating additional boundary conditions, such as Neumann and Robin, are promising future directions. Moreover, although our framework achieves runtimes comparable to existing neural operator baselines, further acceleration of the numerical integration step during the forward pass would facilitate its practical application.

Acknowledgments

This work was supported by the NRF of Korea (RS-2023-00209723); IITP grants (RS-2022-II220594, RS-2023-00227592, RS-2024-00399817, RS-2025-25441313, RS-2025-25443318, RS-2025-02653113); and the Technology Innovation Program (RS-2025-02317326), all funded by the Korean government (MSIT and MOTIE), as well as by the DRB-KAIST SketchTheFuture Research Center.

References

- [1] cholespy: An easily integrable cholesky solver on cpu and gpu. URL <https://github.com/rgl-epfl/cholespy>. Software available on GitHub.
- [2] Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis, A., Dean, J., Devin, M., Ghemawat, S., Goodfellow, I., Harp, A., Irving, G., Isard, M., Jia, Y., Jozefowicz, R., Kaiser, L., Kudlur, M., Levenberg, J., Mané, D., Monga, R., Moore, S., Murray, D., Olah, C., Schuster, M., Shlens, J., Steiner, B., Sutskever, I., Talwar, K., Tucker, P., Vanhoucke, V., Vasudevan, V., Viégas, F., Vinyals, O., Warden, P., Wattenberg, M., Wicke, M., Yu, Y., and Zheng, X. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. URL <https://www.tensorflow.org/>. Software available from tensorflow.org.
- [3] Alkin, B., Fürst, A., Schmid, S., Gruber, L., Holzleitner, M., and Brandstetter, J. Universal physics transformers. In *NeurIPS*, 2024.
- [4] Barill, G., Dickson, N., Schmidt, R., Levin, D. I., and Jacobson, A. Fast winding numbers for soups and clouds. *ACM TOG*, 2018.
- [5] Boullé, N., Nakatsukasa, Y., and Townsend, A. Rational neural networks. In *NeurIPS*, 2020.
- [6] Boullé, N., Earls, C. J., and Townsend, A. Data-driven discovery of green's functions with human-understandable deep learning. *Scientific Reports*, 2022.
- [7] Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., and Zhang, Q. JAX: composable transformations of Python+NumPy programs, 2018. URL <http://github.com/jax-ml/jax>.
- [8] Cao, S. Choose a transformer: fourier or galerkin. In *NeurIPS*, 2021.
- [9] Choromanski, K., Likhoshesterov, V., Dohan, D., Song, X., Gane, A., Sarlos, T., Hawkins, P., Davis, J., Mohiuddin, A., Kaiser, L., Belanger, D., Colwell, L., and Weller, A. Rethinking attention with performers. In *ICLR*, 2021.
- [10] Cooley, J. W. and Tukey, J. W. An algorithm for the machine calculation of complex fourier series. *Math. Comput.*, 1965.
- [11] Deep-MI. Lapy: Toolbox for differential geometry on triangle and tetrahedra meshes, 2023. URL <https://github.com/Deep-MI/LaPy>. Software available on GitHub.
- [12] Hao, Z., Ying, C., Wang, Z., Su, H., Dong, Y., Liu, S., Cheng, Z., Zhu, J., and Song, J. Gnot: A general neural operator transformer for operator learning. In *ICML*, 2023.
- [13] Hu, Y., Zhou, Q., Gao, X., Jacobson, A., Zorin, D., and Panozzo, D. Tetrahedral meshing in the wild. *ACM TOG*, 2018.
- [14] Hu, Y., Schneider, T., Wang, B., Zorin, D., and Panozzo, D. Fast tetrahedral meshing in the wild. *ACM TOG*, 2020.
- [15] Jacobson, A., Panozzo, D., et al. libigl: A simple C++ geometry processing library, 2018. <https://libigl.github.io/>.
- [16] Kim, S., Chi, H.-g., Hu, X., Huang, Q., and Ramani, K. A large-scale annotated mechanical components benchmark for classification and retrieval tasks with deep neural networks. In *ECCV*, 2020.
- [17] Kingma, D. and Ba, J. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- [18] Kitaev, N., Kaiser, , and Levskaya, A. Reformer: The efficient transformer. In *ICLR*, 2020.
- [19] Krishnapriyan, A. S., Gholami, A., Zhe, S., Kirby, R., and Mahoney, M. W. Characterizing possible failure modes in physics-informed neural networks. In *NeurIPS*, 2021.
- [20] Li, Z., Kovachki, N. B., Azizzadenesheli, K., Liu, B., Bhattacharya, K., Stuart, A., and Anandkumar, A. Neural Operator: Graph Kernel Network for Partial Differential Equations. *arXiv*, 2020.
- [21] Li, Z., Kovachki, N. B., Azizzadenesheli, K., Liu, B., Bhattacharya, K., Stuart, A., and Anandkumar, A. Fourier Neural Operator for Parametric Partial Differential Equations. In *ICLR*, 2021.
- [22] Li, Z., Kovachki, N. B., Choy, C., Li, B., Kossaifi, J., Otta, S. P., Nabian, M. A., Stadler, M., Hundt, C., Azizzadenesheli, K., and Anandkumar, A. Geometry-Informed Neural Operator for Large-Scale 3D PDEs. In *NeurIPS*, 2023.

- [23] Li, Z., Meidani, K., and Farimani, A. B. Transformer for partial differential equations' operator learning. *TMLR*, 2023.
- [24] Li, Z., Shu, D., and Farimani, A. B. Scalable transformer for pde surrogate modeling. In *NeurIPS*, 2023.
- [25] Li, Z., Yang, G., Deng, X., De Sa, C., Hariharan, B., and Marschner, S. Neural Caches for Monte Carlo Partial Differential Equation Solvers. In *SIGGRAPH Asia*, 2023.
- [26] Li, Z., Huang, D. Z., Liu, B., and Anandkumar, A. Fourier neural operator with learned deformations for pdes on general geometries. *JMLR*, 2024.
- [27] Nam, H. C., Berner, J., and Anandkumar, A. Solving Poisson Equations Using Neural Walk-on-Spheres. In *ICML*, 2024.
- [28] Negi, P., Cheng, M., Krishnamurthy, M., Ying, W., and Li, S. Learning domain-independent green's function for elliptic partial differential equations. *Computer Methods in Applied Mechanics and Engineering*, 2024.
- [29] Paszke, A., Gross, S., Chintala, S., Chanan, G., Yang, E., DeVito, Z., Lin, Z., Desmaison, A., Antiga, L., and Lerer, A. Automatic differentiation in pytorch. In *NeurIPS Workshop*, 2017.
- [30] Raissi, M., Perdikaris, P., and Karniadakis, G. E. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- [31] Teng, Y., Zhang, X., Wang, Z., and Ju, L. Learning green's functions of linear reaction-diffusion equations with application to fast numerical solver. In *Proceedings of Mathematical and Scientific Machine Learning*, 2022.
- [32] Umetani, N. and Bickel, B. Learning three-dimensional flow for interactive aerodynamic design. *ACM TOG*, 2018.
- [33] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L. u., and Polosukhin, I. Attention is all you need. In *NeurIPS*, 2017.
- [34] Wang, T. and Wang, C. Latent neural operator for solving forward and inverse pde problems. In *NeurIPS*, 2024.
- [35] Wu, H., Luo, H., Wang, H., Wang, J., and Long, M. Transolver: A fast transformer solver for pdes on general geometries. In *ICML*, 2024.
- [36] Xiao, Z., Hao, Z., Lin, B., Deng, Z., and Su, H. Improved operator learning by orthogonal attention. In *NeurIPS*, 2024.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: Yes, the abstract and introduction conveys the core idea of the main text.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [No]

Justification: While we do not discuss the limitation in the main paper, we will include it in the revision.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: The paper does not make theoretical claims that require proofs.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper provides detailed experiment setup required to reproduce the results presented in the paper.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We plan to release the code and data required to reproduce the results upon acceptance.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: The paper discuss the settings of the experiments reported in the main text.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: No, the reported results do not include error bars.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We provide an runtime analysis.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We find no significant ethical issue while working on this work.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [No]

Justification: While the paper does not include statements on broader impacts, we plan to include it in the revision.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: We believe that the models and data used in this work have low potential of misuse.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We have included references to the original datasets used in the experiments.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: We will release the dataset and a documentation upon acceptance.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorosity, or originality of the research, declaration is not required.

Answer: [NA]

Justification: The core method development in this research does not involve LLMs as any important, original, or non-standard components.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

Appendix

In the following, we present the derivation of Eqn. 6 in Sec. A, details of the problem templates used to define the source and boundary functions for the steady-state thermal analysis experiment (Sec. 5.3) in Sec. B, details of the runtime analysis presented in Sec. 5.3 in Sec. C, and additional qualitative results in Sec. D.

A Derivation of the Solution for the Linear System

While the derivation of Eqn. 6 from Eqn. 5 is straightforward, we include the step-by-step derivation for completeness.

Let us decompose the solution \mathbf{u} into interior and boundary components using the section matrices \mathbf{K} and \mathbf{S} :

$$\mathbf{u} = \mathbf{K}^T \mathbf{u}_{\text{int}} + \mathbf{S}^T \mathbf{h}, \quad (15)$$

where $\mathbf{u}_{\text{int}} \in \mathbb{R}^{(N_v - N_b)}$ denotes the solution of Eqn. 5 within the domain.

As \mathbf{u} satisfies Eqn. 5, we have:

$$\mathbf{L} (\mathbf{K}^T \mathbf{u}_{\text{int}} + \mathbf{S}^T \mathbf{h}) = \mathbf{Mf}. \quad (16)$$

Expanding the left-hand side and rearranging terms yields:

$$\mathbf{L} \mathbf{K}^T \mathbf{u}_{\text{int}} = \mathbf{Mf} - \mathbf{L} \mathbf{S}^T \mathbf{h}. \quad (17)$$

By left-multiplying \mathbf{K} , we obtain

$$\mathbf{K} \mathbf{L} \mathbf{K}^T \mathbf{u}_{\text{int}} = \mathbf{K} \mathbf{Mf} - \mathbf{K} \mathbf{L} \mathbf{S}^T \mathbf{h}. \quad (18)$$

Assuming $\mathbf{K} \mathbf{L} \mathbf{K}^T$ is full-rank, which is ensured by imposing appropriate boundary conditions, \mathbf{u}_{int} can be written as:

$$\mathbf{u}_{\text{int}} = \mathbf{G} (\mathbf{K} \mathbf{Mf} - \mathbf{K} \mathbf{L} \mathbf{S}^T \mathbf{h}), \quad (19)$$

where $\mathbf{G} = (\mathbf{K} \mathbf{L} \mathbf{K}^T)^{-1}$. Substituting this result into Eqn. 15 gives us Eqn. 6.

B Details of Problem Templates

When constructing the benchmark used for the steady-state thermal analysis experiment in Sec. 5.3, the source functions f are populated from the following template:

$$\begin{aligned} f(x, y, z) = & \Delta \{ \sin A\pi x \cdot \cos AC\pi y \\ & + (1 - \cos A\pi x) \cdot (1 - \sin AB\pi y) \\ & + \sin^2 AD\pi z \}, \end{aligned} \quad (20)$$

with the coefficients $A = \{1.25, 2.5\}$, $B = \{1.5, 3.5\}$, $C = \{1.5, 3.5\}$, and $D = \{1.5, 3.5\}$, resulting in 16 different combinations. Similarly, the boundary functions are generated using the template:

$$\begin{aligned} h(x, y, z) = & E(x^3 - 3xy^2) + F(y^3 - 3x^2y) \\ & + (x^2 - z^2), \end{aligned} \quad (21)$$

using 4 distinct combinations of the coefficients: $E = \{-1.0, 1.0\}$ and $F = \{0.0, 1.0\}$.

C Details of Runtime Analysis

In this section, we provide additional details of runtime analysis in Sec. 5.3. We use the binary compiled from the official implementation of `fTetWild` [14] for generating tetrahedral meshes. The implementation supports multi-core processing, and we used the default parameters with an edge length of 0.02 for mesh generation. To populate the query points used for measuring runtimes, we

	Meshing (s)	Solve (s)	FEM Total (s)	IPQ (s)	Ours Forward (s)	Ours Total (s)
SCREWS & BOLTS	12.579	0.377	12.956	0.016	0.023	0.039
NUT	11.601	0.637	12.238	0.014	0.038	0.051
MOTOR	48.661	1.434	50.095	0.082	0.058	0.140
FITTING	18.037	0.401	18.439	0.024	0.035	0.059
GEAR	45.803	0.581	46.384	0.180	0.044	0.225

Table 6: Breakdown of runtime comparison with the FEM solver. The total runtime of the FEM solver includes both the meshing and linear solve times, whereas our runtime is computed as the sum of the time required for the interior point query (denoted IPQ) and the network forward pass. All runtimes are reported in seconds.

	Transolver (Wu et al. [35])			LNO (Wang & Wang [34])		UPT (Alkin et al. [3])		Ours	
	IPQ (s)	Forward (s)	Total (s)	Forward (s)	Total (s)	Forward (s)	Total (s)	Forward (s)	Total (s)
SCREWS & BOLTS	0.016	0.017	0.033	0.006	0.022	0.040	0.056	0.023	0.039
NUT	0.014	0.008	0.022	0.006	0.020	0.062	0.076	0.038	0.051
MOTOR	0.082	0.008	0.090	0.006	0.088	0.084	0.166	0.058	0.140
FITTING	0.024	0.008	0.032	0.006	0.030	0.052	0.076	0.035	0.059
GEAR	0.180	0.008	0.188	0.007	0.187	0.066	0.246	0.044	0.225

Table 7: Breakdown of runtime comparison with neural operators. The total runtime of each neural operator is computed as the sum of the time required for the interior point query (denoted IPQ) and the network forward pass. All runtimes are reported in seconds.

uniformly sample the 3D space and determine whether each point lies inside the domain using the Fast Winding Number [4] implemented in `libig1` [15]. To ensure a fair comparison with our method, which leverages GPU acceleration during network inference, we re-implemented the solver algorithm from [11] to support GPU acceleration. In particular, we utilize `Cholespy`[1], a GPU-accelerated Cholesky solver, for matrix prefactorization to solve Eqn. 5. The runtime breakdowns corresponding to the results in Tab. 3 are presented in Tab. 6 and 7, respectively, each of which summarizing the time spent in each stage of the FEM solver—including meshing and solving linear systems—as well as, for the neural operators, the time required for interior point queries (IPQ) and network forward passes. All runtimes are measured on a system equipped with an Intel Xeon Gold 6442Y processor with 24 cores and an NVIDIA RTX 3090 GPU with 24 GB of VRAM.

D Additional Qualitative Results

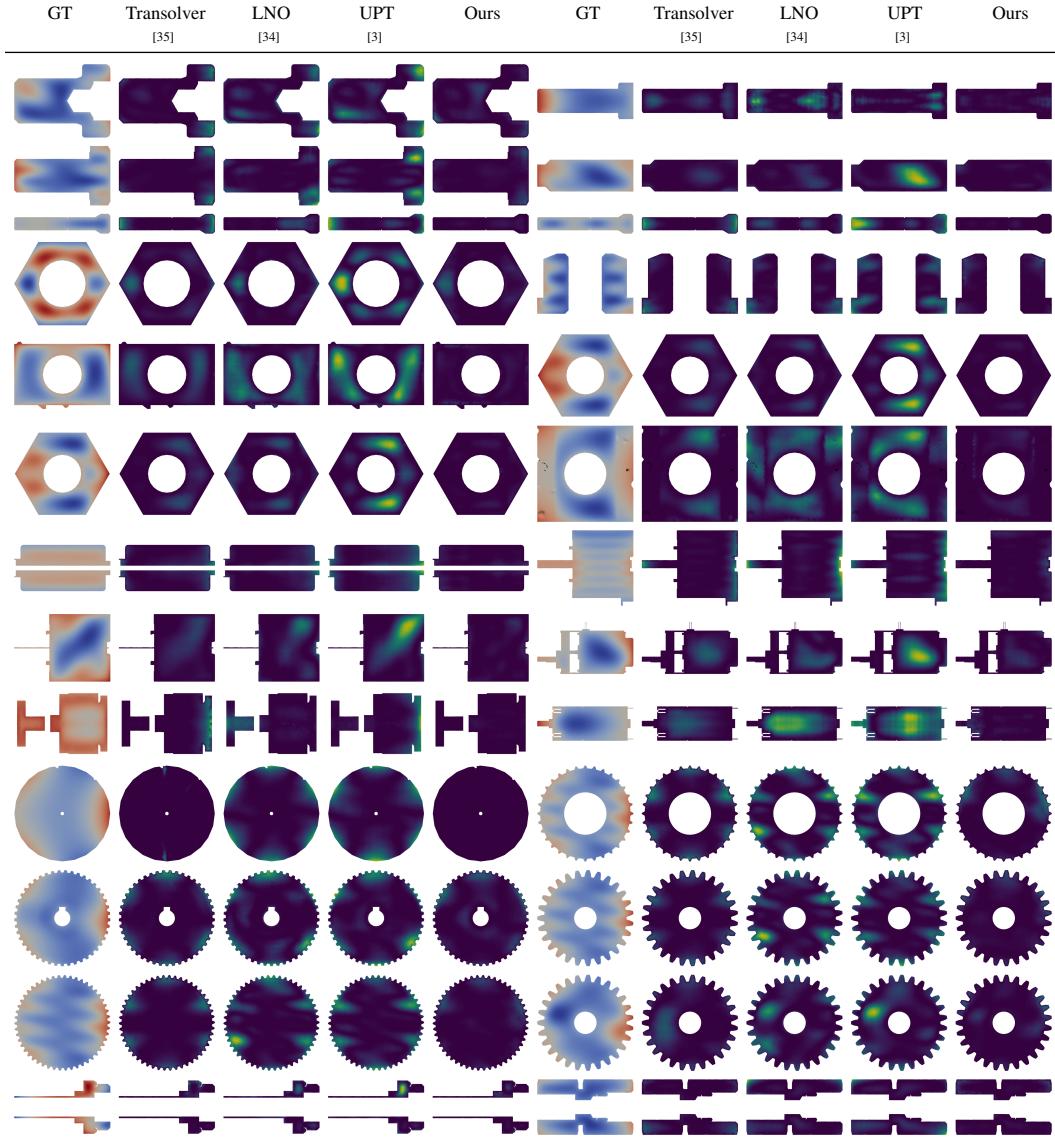


Figure 6: Additional qualitative comparison. For each ground-truth solution in columns one and six, the per-point L_2 errors from various methods are visualized in the error maps across the remaining columns.